

The Missing Metrics in DSA Content Moderation Transparency

January 8, 2026 / in Analysis

By Max Davy, Oxford Internet Institute

The Digital Services Act makes platform transparency reporting mandatory and standardised, but the metrics it requires still fall short of what is needed for real accountability. Counts of removals and appeals alone cannot tell us whether content moderation systems are accurate, proportionate, or effective, making the absence of evaluation metrics such as precision and recall increasingly difficult to justify under the DSA's risk-based logic. While these metrics are unlikely to surface in baseline transparency reports under Articles 15 and 24, the post argues they may yet emerge through heightened scrutiny of the largest online platforms and search engines (VLOPSEs), as regulatory expectations take shape through enforcement, systemic risk reporting, audits, and related obligations.

How much do you like stumbling across porn on Instagram? Me neither – and as we might expect, it violates Meta's policies. But not all content

Recent posts

If at first you don't succeed: Reflections on a rejected Art. 40 DSA data access request March 12, 2026

How Have Platforms Addressed Addictive Design Under DSA March 9, 2026

Shareholder Control and the New Politics of Platform Regulation March 6, 2026

How has the DSA performed in protecting election integrity? March 4, 2026

Reclaiming the Algorithm: A



[<https://transparency.meta.com/policies/community>

standards/?source=https%3A%2F%2Fen-

gb.facebook.com%2Fcommunitystandards%2F]

(think of Meta's complicated history

[<https://www.theguardian.com/technology/2023/jan/17/fre>

the-nipple-meta-facebook-instagram] with

censoring nipples).

If we want to assess whether Meta effectively

moderates the many millions of posts that may

violate its policies in this area, we can start by

looking at how much content Meta actions over

time, courtesy of their Community Standards

Enforcement Report

[[https://transparency.meta.com/reports/community-](https://transparency.meta.com/reports/community-standards-enforcement/adult-nudity-and-sexual-activity/facebook/)

standards-enforcement/adult-nudity-and-sexual-

activity/facebook/] :



Notice the peak in late 2024, when the number of

Meta's enforcement actions (66.6 million) against

adult nudity and sexual activity content roughly

doubled compared to the previous quarter?

2020

Who speaks and

who is heard?

Civil society

participation

and

participatory

justice in DSA

systemic risk

management

February 17,

2026

Reclaiming the

Algorithm:

What the DSA

can—and can't—

fix about

recommender

systems

February 15,

2026

The Missing

Metrics in DSA

Content

Moderation

Transparency

January 8,

2026

What are DSA

audits doing for

systemic risk

enforcement?


The case of X



content moderation system became more effective at catching it. More concerning would be if Meta's system had started over-blocking material which *didn't* violate its Community Standards (whether because of changes to company policies, changes to its automated systems, or any number of other reasons).

If this were true, and the spike in enforcement actions happened because Meta was systematically classifying more innocent material as violative, we might then expect more users to complain to the platform and appeal enforcement actions.

Thanks to Meta's transparency report, we can then check whether this was the case by looking at the total number of user appeals:




Based on this chart, it appears the total number of appeals during this period stayed level. But we should also want to know the success rate of those

Enforcement

Moment

November 26,
2025

ALL
POSTS



So, although the total number of appeals stayed level, this chart shows that the number of *successful* appeals more than doubled. This suggests that, when Meta ramped up its enforcement actions against adult nudity and sexual activity content in late 2024, there was at least some corresponding increase in content being unjustly blocked or actioned against.

It's tempting to estimate, based on the number of successful appeals, the **precision** of Meta's enforcement actions: that is, the percentage of content Meta legitimately (and illegitimately) blocks under its community guidelines. However, not every innocent person whose content is blocked on Facebook appeals the verdict – and the propensity to appeal varies considerably from person to person and group to group [<https://www-tandfonline-com.ezproxy-prd.bodleian.ox.ac.uk/doi/full/10.1080/1369118X.2024.23>]



really? Meta could at least help resolve the ambiguity by releasing their own direct estimates of the precision of their enforcement actions. But they don't, and so we cannot. Our investigatory journey – at least the one which relies on Meta's transparency disclosures – ends.

There are two lessons to be drawn from this.

Firstly, when trying to understand an automated system, it is never sufficient to analyse metrics which are straightforward counts of the digital traces it leaves behind. Automated systems, by definition, do not believe that they make mistakes. Evaluative human judgement is needed for any metric aiming to produce evaluative information about a system.

Secondly, our analysis thus far relies on data that Meta provides to us. Such disclosures, once released voluntarily, are now obligatory under the *Digital Services Act* (DSA). Articles 15 and 24 of the DSA oblige providers of intermediary services and online platforms, respectively, to periodically release transparency reports describing various aspects of their content moderation practices. This is a unique mechanism to force companies to publish metrics regardless of whether it's good for their reputation.



Evaluation vs operational metrics

It's useful to draw a distinction between two types of metrics.

Operational metrics are those which describe the operations of a digital platform – the number of posts, blocks, messages, appeals, complaints, and so on. These are important data, but as we saw in the introductory example, raw data (e.g. # of blocks, # of appeals, etc) are insufficient for evaluating the quality of a moderation system. This is where evaluation metrics become necessary.

Evaluation metrics are metrics augmented by human judgements about whether enforcement actions are legitimate in accordance with a 'ground truth', i.e. the platform's own moderation rules and policies. Examples include *precision* (what % of blocked content was actually violative?), *recall* (what % of violative content was caught?) and *prevalence* (what % of the (viewed) content on a platform is violative?).

The insufficiency of operational metrics is roughly analogous to the 'dark figure' in crime statistics [<https://www.ojp.gov/ncjrs/virtual-library/abstracts/understanding-crime-data-haunted-dark-figure>] , which refers to crimes which go unreported and so remain invisible in official



produce additional evaluative information. In

contrast, because tech companies have complete observability over activity on their platforms, they are able to sample content or activity and make these evaluative judgements themselves, rather than having to triangulate from multiple unreliable metrics.

We know from trust and safety professionals, including from Google

[<https://www.unofficialgoogledatascience.com/2019/08/es-prevalence-of-rare-events.html>] and Facebook

[<https://www.youtube.com/watch?v=v-kPwynEYE4>] , that evaluation metrics are critically important to internal decision-making in tech companies, and this mirrors my experience as a former data scientist in Canva's Trust and Safety team. Here's a quote

[<https://www.tsipa.org/curriculum/ts-fundamentals/transparency-report/transparency-report-categories/>] from the Trust and Safety Professional Association's curriculum: "Metrics like these can be challenging to generate, but act as an important accountability mechanism that incentivizes companies to continuously improve the true effectiveness of their policies and enforcement systems."



transparency reports

I'm not the first person to point out this problem. In 2019, a group of independent experts released a report

[<https://law.yale.edu/sites/default/files/area/center/justice/c> commissioned by Meta (then Facebook) evaluating the metrics in its transparency reports. They called on Facebook to publish measures of recall and precision and similar evaluation metrics for both human and automated judgements. One quote is directly relevant to our introductory example: “Rates of reversal on appeal should be made public, but they should not stand in as the sole public metric of accuracy.”

Facebook then published a non-committal response [<https://research.facebook.com/blog/2019/5/exploring-feedback-from-data-and-governance-experts-a-research-based-response-to-the-data-transparency-advisory-group-report/>] in which it argued that ambiguous moderation cases would complicate evaluative metrics – despite the fact that the company already used such metrics internally. It's been more than six years since this report was released, and the suggested evaluation metrics still haven't been published externally.



expand and maintain their power by gathering their sources of information about themselves without sharing it publicly or with regulators. Although platforms might protest otherwise, their reluctance to engage in voluntary transparency and facilitate outsider scrutiny is consistent with recent reductions in Trust & Safety investments [<https://dl-acm-org.ezproxy-prd.bodleian.ox.ac.uk/doi/10.1145/3706598.3713662>] more generally.

We know from leaked internal Meta documents [<https://www.reuters.com/investigations/meta-is-earning-fortune-deluge-fraudulent-ads-documents-show-2025-11-06/>] that Meta estimated 10% its advertisement revenue came from scammers. This is a highly consequential evaluation metric, but was never intended for public consumption. If companies were compelled to publish such metrics themselves, they could find it more difficult to muddy the public debate by claiming, as Meta did in this case, that the leaked evidence “present[s] a selective view that distorts Meta’s approach to fraud and scams.”

The DSA was an opportunity to force companies such as Meta to publish such metrics. The concepts measured by evaluation metrics cleanly map onto [<https://dl.acm.org/doi/10.5555/3716662.3716795>]



communicates the extent to which users' freedom of expression is being curtailed. These metrics also aid in understanding how companies are balancing between these rights.

Yet the DSA does not call for a single evaluation metric – with one exception. Article 15.1(e) requires platforms to publish 'indicators of accuracy' of automated moderation. However, this requirement is both limited in its scope of 'automated systems' and has many other problems, so that the information provided under this provision thus far are essentially meaningless.

Any Trust & Safety data scientist could tell you that evaluation metrics such as recall and precision should be included for any 'transparency report' worth its name. It begs the question: how did the EU choose what metrics to include?

How metrics were chosen for the DSA

If transparency programs are enacted without reference to a goal, they become useless

[https://scholarship.law.columbia.edu/faculty_scholarship/

. The design of any transparency report should start by asking the intended purpose of providing the information, use this purpose to determine what information ought to be conveyed, and only then design metrics to address those aims.



[[https://ec.europa.eu/info/law/better-](https://ec.europa.eu/info/law/better-regulation/have-your-say/initiatives/12417-Digital-Services-Act-deepening-the-Internal-Market-and-clarifying-responsibilities-for-digital-services/public-consultation_en)

[regulation/have-your-say/initiatives/12417-Digital-Services-Act-deepening-the-Internal-Market-and-clarifying-responsibilities-for-digital-services/public-consultation_en](https://ec.europa.eu/info/law/better-regulation/have-your-say/initiatives/12417-Digital-Services-Act-deepening-the-Internal-Market-and-clarifying-responsibilities-for-digital-services/public-consultation_en)] (from June to September 2020) suggests they did not. In the Commission’s questionnaire, one question asked companies that already published voluntary transparency reports what metrics were included in those reports:

3 Do the reports include information on:

- ☒ X Volumes of takedowns and account suspensions following enforcement of your terms of service?
- ☒ X Volumes of takedowns following a legality assessment?
- ☒ X Notices received from third parties?
- ☐ Referrals from authorities for violations of your terms of service?
- ☐ Removal requests from authorities for illegal activities?
- ☒ X Volumes of complaints against removal decisions?
- ☒ X Volumes of reinstated content?
- ☐ Other, please specify in the text box below

Screenshot from the Commission’s 2020 consultation questionnaire

This list is *very* similar to the final set of content moderation metrics stipulated for disclosure under the DSA’s transparency obligations.

Some companies (including, funnily enough, Facebook [https://ec.europa.eu/info/law/better-regulation/have-your-say/initiatives/12417-Digital-Services-Act-deepening-the-Internal-Market-and-clarifying-responsibilities-for-digital-services/public-consultation_en]) did recommend, in their consultation response, alternate methods of choosing metrics. Nevertheless, it appears there



transparency reporting metrics already being performed by Facebook, Google, and other firms in the DSA.

To summarise our analysis: Evaluation metrics clearly align with the DSA's purpose. However, despite being requested by previous expert panels, despite companies already calculating such metrics for internal use, and despite industry suggestions that the proposed metrics were insufficient, the DSA transparency reporting legislation did not include them.

Can we do anything about it?

Inflexible legislation

One surprising aspect of how the DSA handles transparency reporting is that it specifies a set of metrics, in prescriptive detail, directly in the legislation. Because changing legislation is difficult and slow, this is a nightmare scenario for a data scientist.

It's commonly understood in industry that metrics need to be quickly iterated upon as edge cases and unexpected subtleties undermine the construct validity [<https://www-tandfonline-com.ezproxy-prd.bodleian.ox.ac.uk/doi/full/10.1080/1369118X.2024.23> of the original definition of metric. This is especially the case for not-yet-well-understood



[<https://www.cambridge.org/core/books/dsa-observatory>]

online-controlled-

experiments/D97B26382EB0EB2DC2019A7A7B518F59

for more on this point.

It seems the DSA does not offer a direct path to iterating on its basic transparency reporting requirements under Articles 15 and 24 (unlike some other provisions, there was no delegated legislation provided to further specify substantive requirements; an implementing act

[https://ec.europa.eu/info/law/better-regulation/have-your-say/initiatives/14027-Digital-Services-Act-transparency-reports-detailed-rules-and-templates_en] *was* used for prescribing the required format of transparency reporting, but did not touch on metric definitions).

Possible alternative routes for requiring evaluation metrics

For now, it appears that the DSA's transparency reporting requirements for companies under Articles 15 and 24 are largely limited to disclosing operational metrics. However, other DSA transparency mechanisms might offer a pathway for getting at least the largest online platforms and search engines (VLOPSEs) to publish evaluation metrics.



transparency requirements under Article 12. Both

the DSA Civil Society Coordination Group

[<https://cdt.org/wp-content/uploads/2025/03/RA-Report-Assessment-Report.pdf>] (pages 15-22) and

the Integrity Institute

[<https://drive.google.com/file/d/1MJHx4cx24XV4UZUfW>

(page 35) have called for corporations to

accompany their claims about the effectiveness of their content moderation systems with actual proof, including quantitative evaluation metrics. The data access mechanism provided by Article 40 of the DSA, too, provides a complementary route for regulators and independent vetted researchers to potentially scrutinize some of the evaluation metrics that VLOPSEs use internally to measure the efficacy of their content moderation systems.

But we're not there yet. The initial independent audit reports [<https://dsa-observatory.eu/2025/12/19/dsa-audits-what-are-they-good-for/>] of VLOPSEs' risk management obligations and the European Commission's summary report [<https://digital-strategy.ec.europa.eu/en/news/digital-services-act-report-lays-out-landscape-systemic-risks-online>] of prominent and recurring risks on VLOPSEs are first steps toward cataloguing risk management practices, but these documents don't yet really engage with how the effectiveness of moderation



[<https://www.channelformerlyand>

professionals/online-safety-

framework/certifications-schemes/vetted-

researchers/] , it will take some time before we

know what kinds of insights can be gleaned from

such access requests.

Conclusion

This post has argued that metrics included in transparency reports must be developed via a process which carefully considers what evaluative concepts ought to be conveyed, and how such concepts can best be operationalised. I have further argued that such a process wasn't followed in the design of the Digital Services Act's transparency reporting framework. The result is a set of metrics that reveal little about the choices, investment, quality, or trade-offs in online platforms' content moderation systems. Other jurisdictions considering similar legislation, and future reform efforts for the DSA, should consider the simple and impactful addition of evaluation metrics to transparency reporting requirements.

In an ideal world, companies would be compelled to publish high-level evaluation metrics for their content moderation practices in a standardised manner. For now, this gap can be partially closed for the largest platforms via Article 42's heightened

At present, platforms’ systemic risk reports and related audit reports offer some detail about companies’ safety efforts, but rarely include evaluative metrics that would allow outsiders to assess whether those measures actually work. Where companies claim their content moderation measures are effective or proportionate, they should be expected to provide evidence in the form of evaluation metrics to back it up.

Share this entry

